

Transformers

Plan

- **Transformers**
 - Multi-Headed Self-Attention
 - Transformer Layers
 - GPT-style Decoders

Self-Attention


- Suppose there are N input tokens $x_1, \dots, x_N \in \mathbb{R}^d$ (e.g. words).
- Self-attention maps the input sequence into another sequence of the same length.
- For each token x_i , we compute a **distribution** $\alpha_i \in \mathbb{R}_+^d$ over all input tokens, given by

$$\alpha_{i,j} = \frac{\exp(x_i \cdot x_j)}{\sum_{j=1}^N \exp(x_i \cdot x_j)}$$

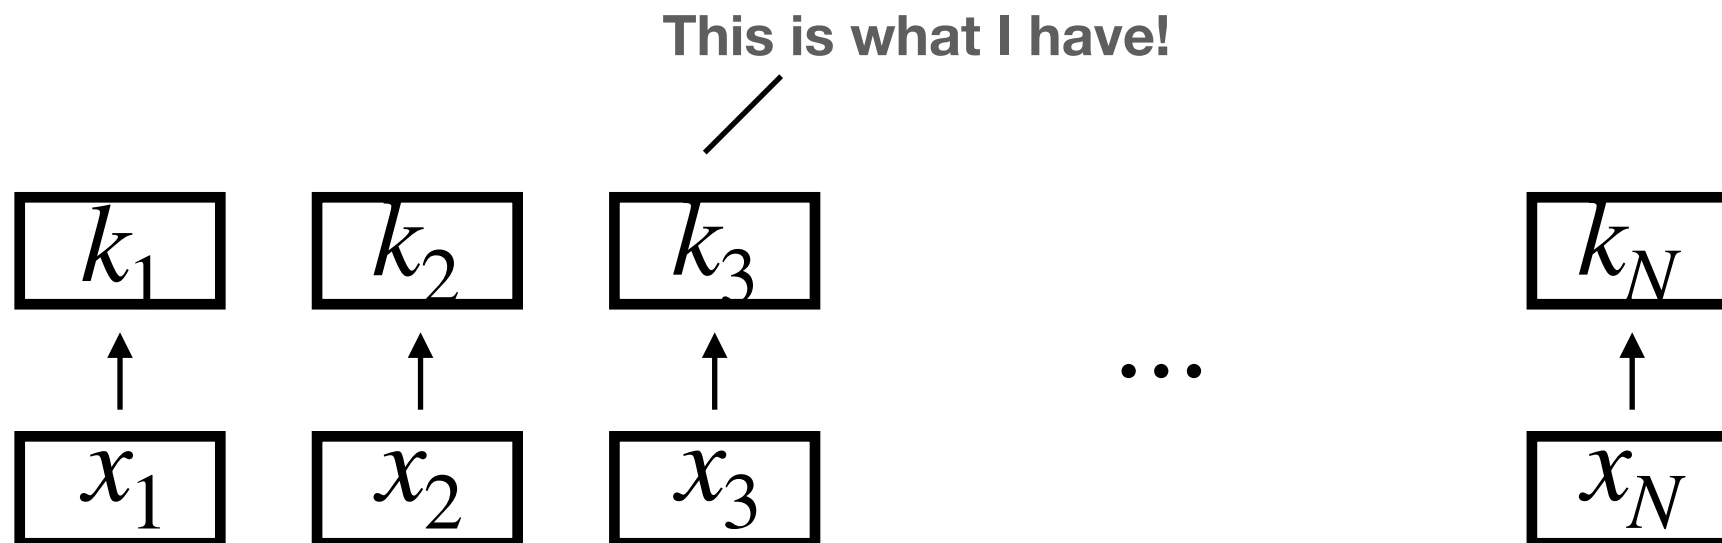
- To compute the i th output, we combine the input vectors using the attention weights α_i .
- We give ourselves more flexibility using **queries** and **keys** to compute the attention weights, and by computing a combination of **values** instead of the original input.

$$q_i = M_q x_i + \beta_q$$

$$k_i = M_k x_i + \beta_k$$



This is what I need!



This is what I have!

o_1

o_2

o_3

o_N

x_1

x_2

x_3

x_N



...



q_1

q_2

q_3

q_N

k_1

k_2

k_3

k_N



...

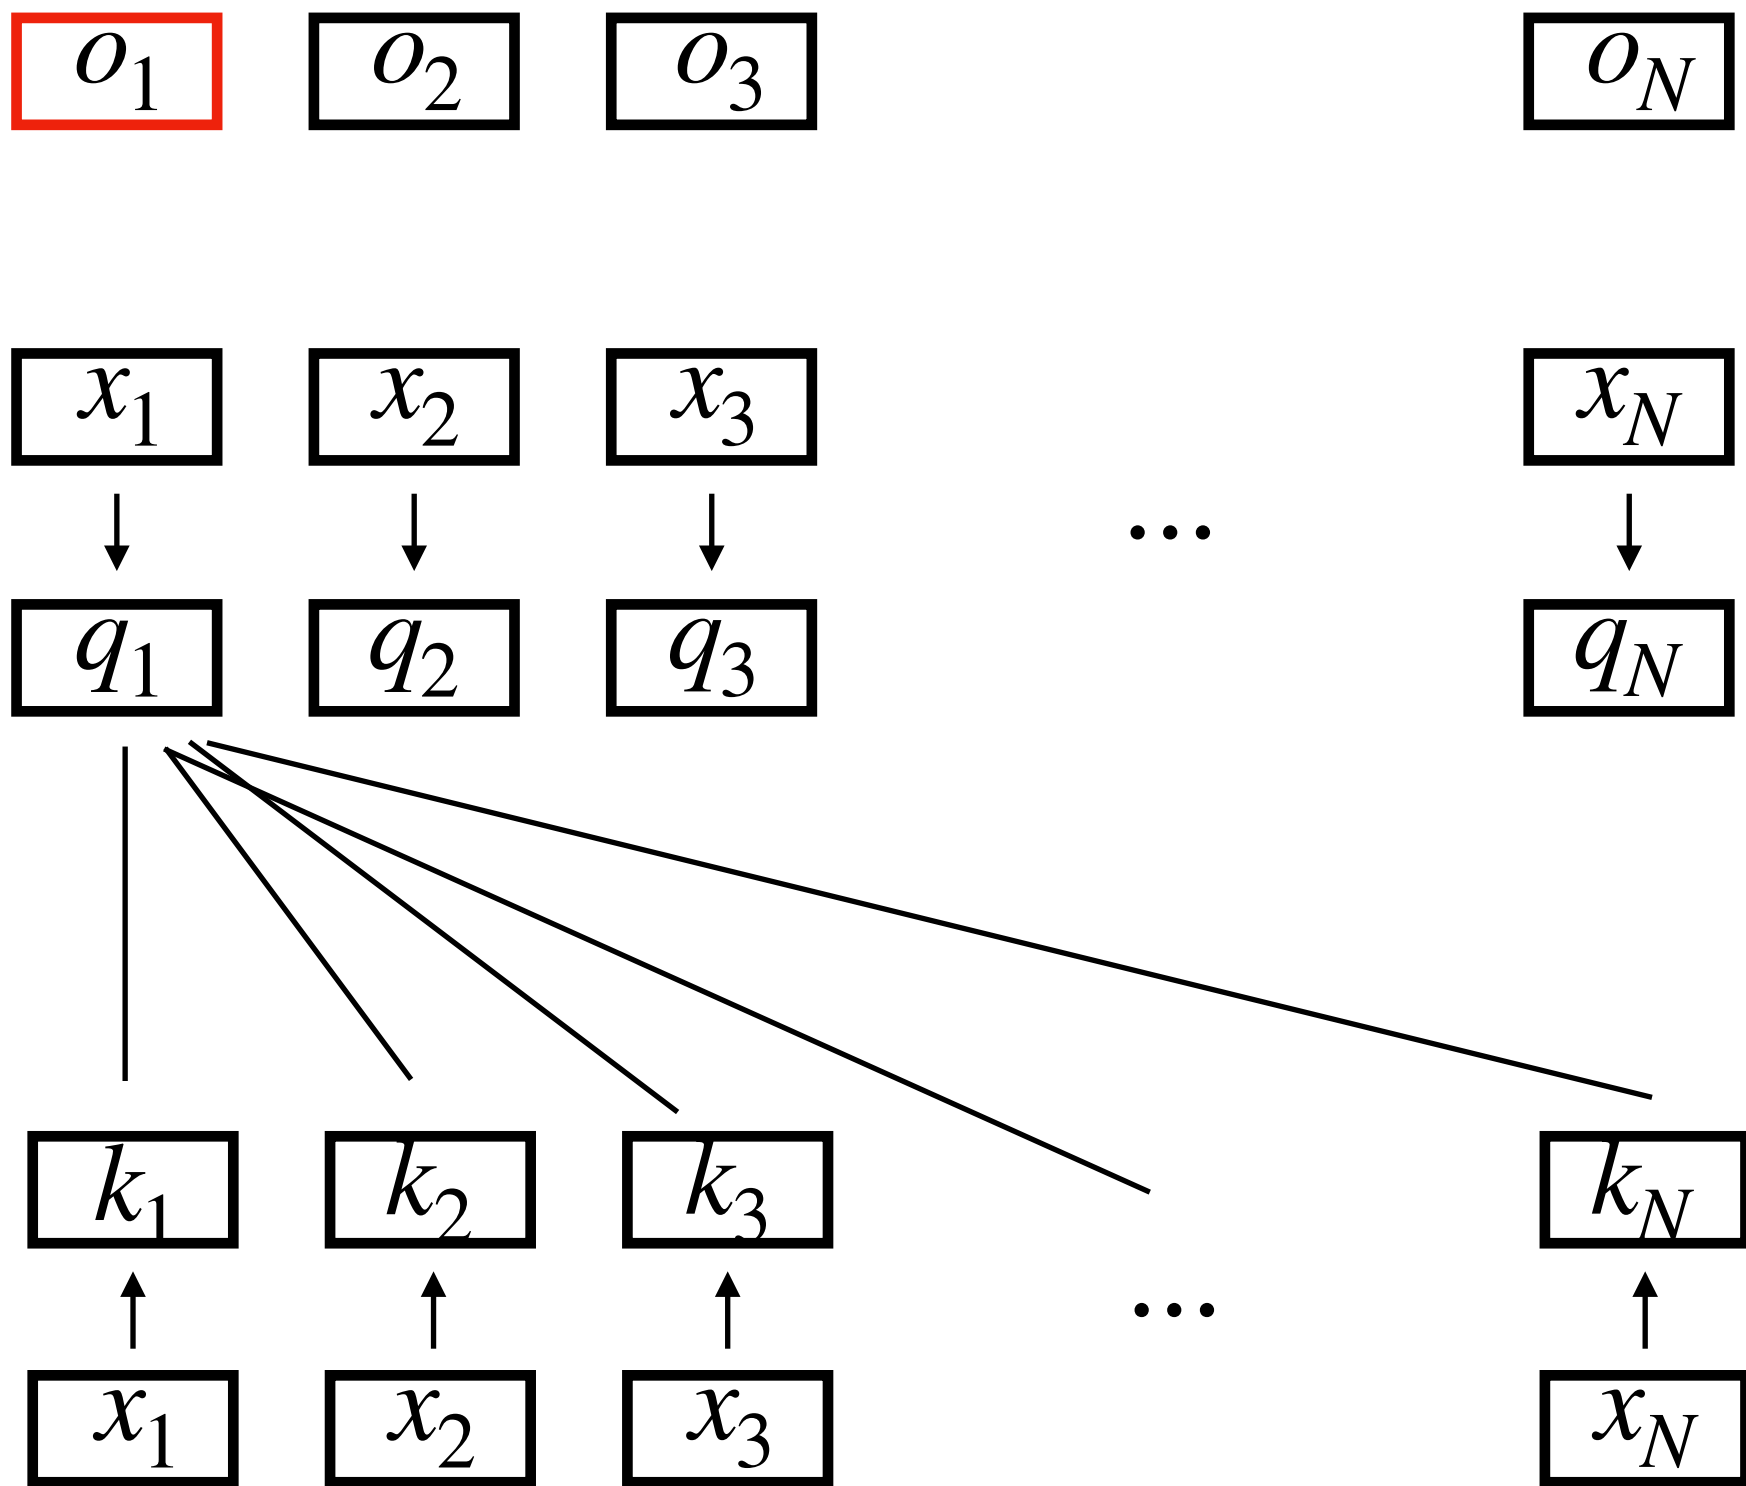


x_1

x_2

x_3

x_N



o_1 o_2 o_3 ... o_N

x_1 x_2 x_3 ... x_N



...



q_1 q_2 q_3 ... q_N

$$\alpha_{1,j} \propto \exp(q_1 \cdot k_j)$$

k_1

k_2

k_3

...

k_N



x_1

x_2

x_3

x_N

$$o_1$$

$$o_2$$

$$o_3$$

$$o_N$$

$$\alpha_{1,j} \propto \exp(q_1 \cdot k_j)$$

$$o_1 = \sum_{i=1}^N \alpha_{1,j} v_j$$

$$v_1$$

$$v_2$$

$$v_3$$

$$v_N$$

$$\uparrow$$

$$x_1$$

$$\uparrow$$

$$x_2$$

$$\uparrow$$

$$x_3$$

...

$$\uparrow$$

$$x_N$$

$$v_i = M_v x_i + \beta_v$$

Scaled Dot-Product Self-Attention

- Arrange the input into a $N \times D$ matrix X (each row is an element of the sequence.)
- Matrices of **queries** $Q = XM_q + \beta_q$, **keys** $K = XM_k + \beta_k$, and **values** $V = XM_v + \beta_v$. Each row is a (query / key / value).

$$\text{SA}(X) = \frac{\text{softmax}(QK^T)}{\sqrt{D}}V$$

(Multi-Head) Scaled Dot-Product Self-Attention

- We have H attention mechanisms (each with its own parameters), of dimension D/H .
- Compute output from each and concatenate.
- Project once more to form the output.

```

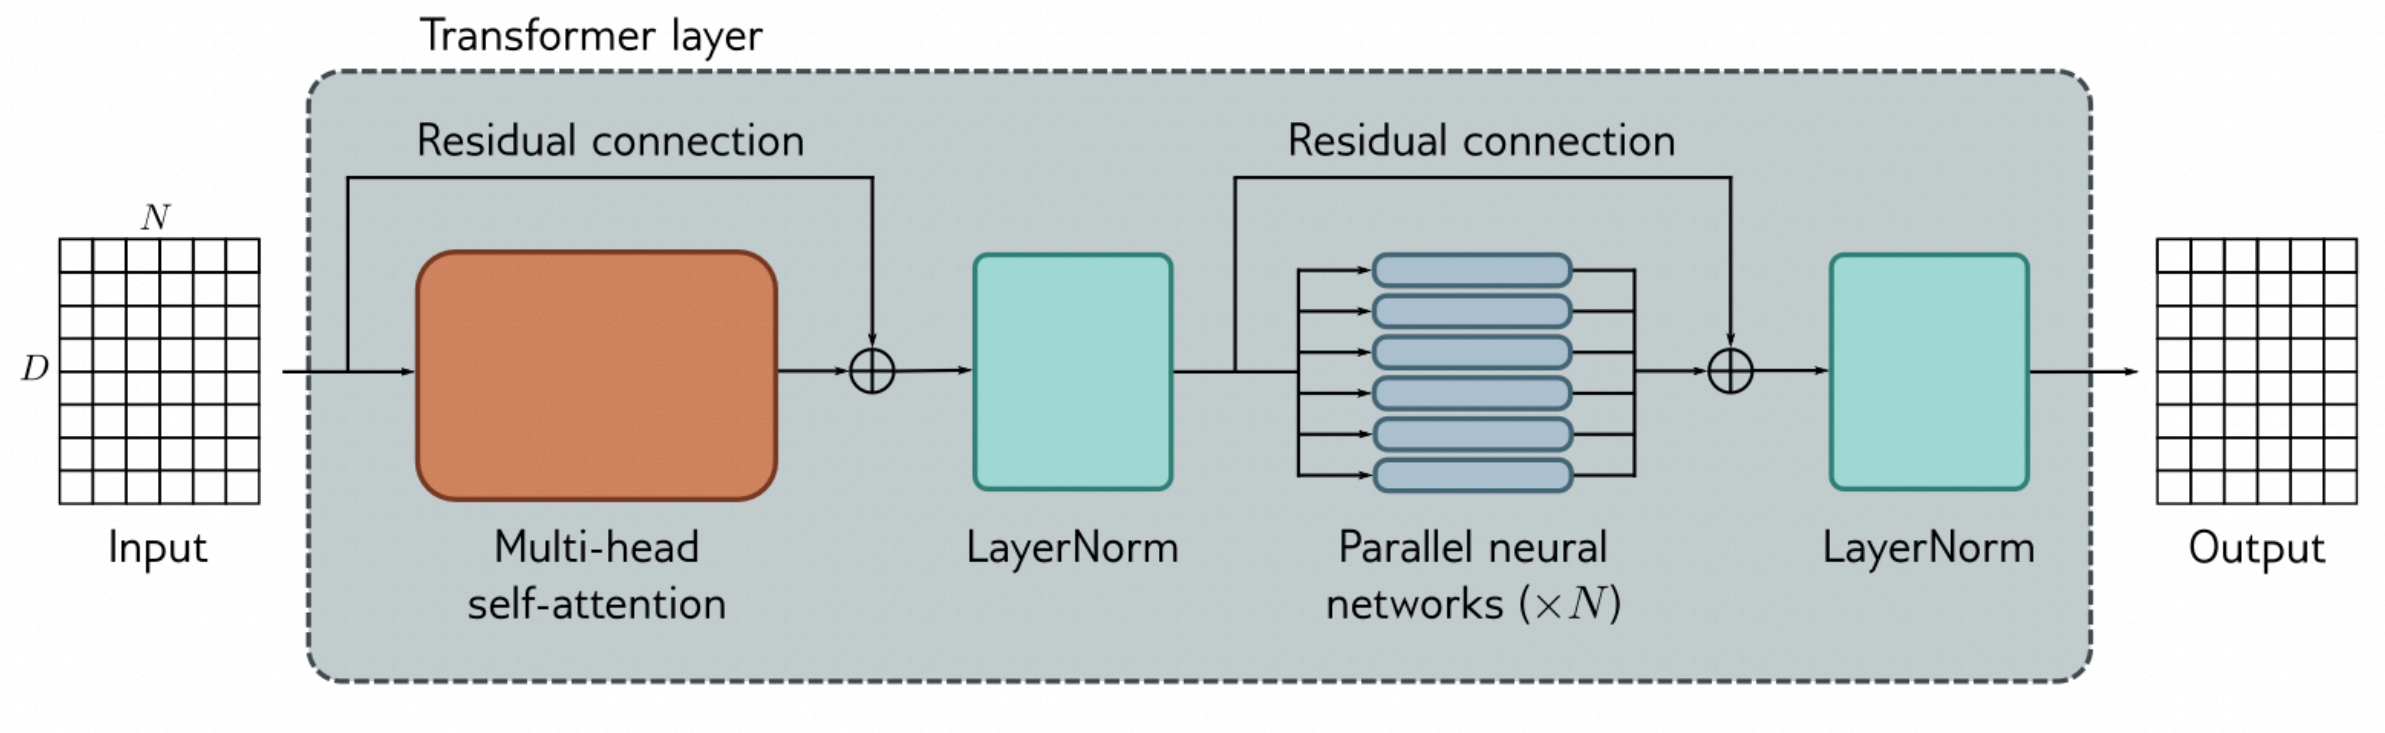
# Compute queries, keys, and values for all heads at once.
q = jax.vmap(self.lin_q)(x).reshape(N, D // n_head, n_head)
k = jax.vmap(self.lin_k)(x).reshape(N, D // n_head, n_head)
v = jax.vmap(self.lin_v)(x).reshape(N, D // n_head, n_head)

# Attends over the values to produce the output values.
# The attention coefficients are masked using jnp.tril.
def sa(q, k, v):
    mask = jnp.triu(jnp.ones((N, N)) * -float('inf'), k=1)
    return jax.nn.softmax(q @ k.T / jnp.sqrt(D // n_head) + mask) @ v

```

Transformer Layers


Transformer Layer



- Actually: **LayerNorm** \rightarrow Self-Attention \rightarrow **LayerNorm** \rightarrow MLP
- On Layer Normalization In the Transformer Architecture [<https://arxiv.org/pdf/2002.04745.pdf>]


Batch Normalization

- Recall that in regression, we predict a **response** y from **covariates** (features) x_1, \dots, x_p .
- What happens if the data we train on and the data we test on are different? One type of such “distribution shift” is **Covariate Shift**.
- The marginal distribution of the covariates is $q(x)$ at train time and $p(x)$ at test time. Here, $p(y | x) = q(y | x)$ - only the marginal of the covariates changes.



Batch Normalization

- **Interval Covariate Shift:** covariate shift is happening **inside** of the network.




Batch Normalization

- Solution: shift and scale the input to each layer.

$$\hat{x}_i^{(k)} = \frac{x_i^{(k)} - \mu_B^{(k)}}{\sqrt{(\sigma_B^{(k)})^2 + \epsilon}}$$



$$y_i^{(k)} = \gamma^{(k)} \hat{x}_i^{(k)} + \beta^{(k)}$$

- Gamma(s) and beta(s) are learned parameters used to ensure that “the transformation inserted in the network can represent the identity transform.”



Layer Normalization


- In BatchNorm, we standardize over each **feature** separately, **across the batch**.
- In LayerNorm, we standardize over each **batch example** separately, across the **features**.



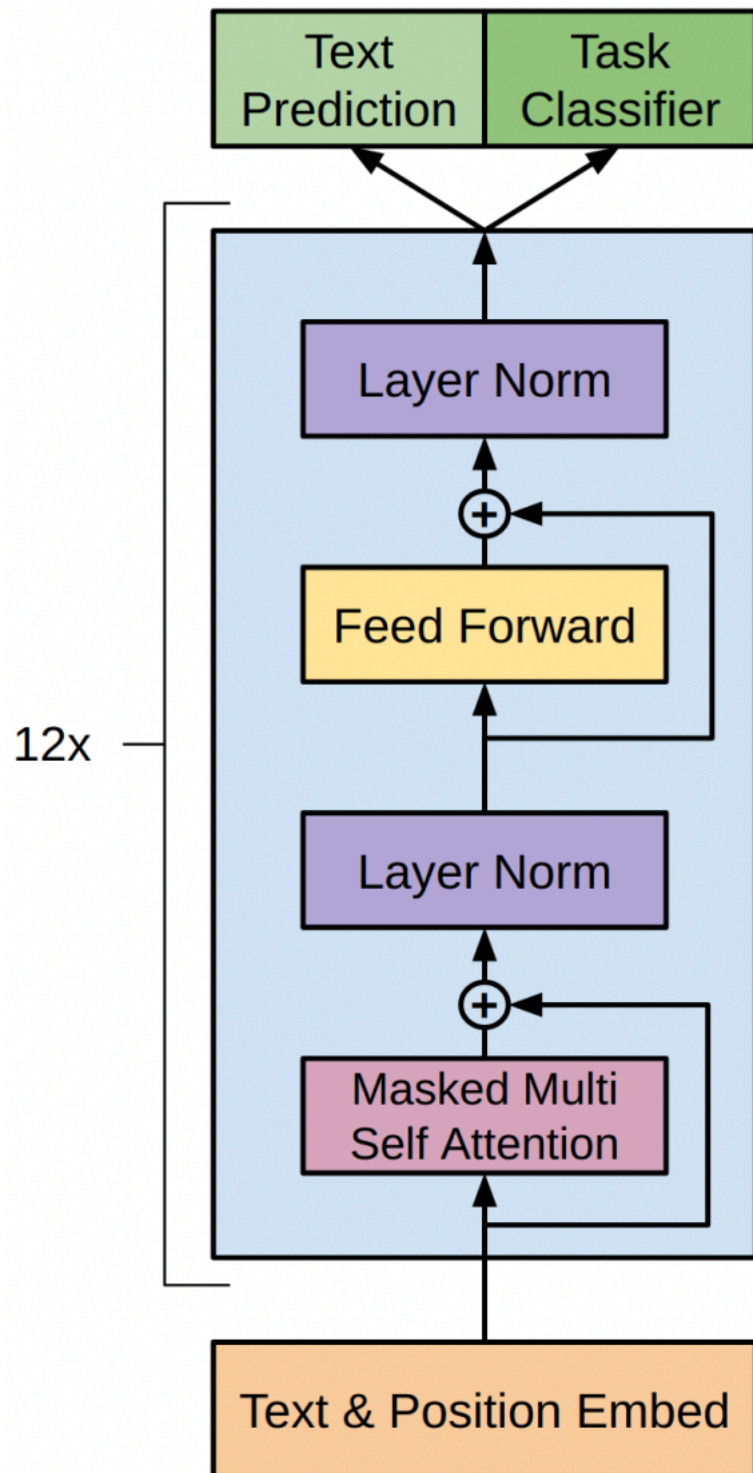
Multilayer Perceptron

- Just a FFN applied to each time-step
 - GPT2: two-layer, with hidden size $4 \cdot D$.
- GeLU (Gaussian Error Linear Unit) activation [<https://arxiv.org/pdf/1606.08415.pdf>]

$$\text{GeLU}(x) = x\Phi(x) \approx 0.5x \left(1 + \tanh \left(\sqrt{\frac{2}{\pi}} (x + 0.044715x^3) \right) \right)$$




GPT2-Style Decoder



GPT2-Style Decoder

- Input string: “attention is all you need”.
- “attention is all you need” —> [1078, 1463, 318, 477, 345, 761] (indices)
- **Word Token Embeddings (WTE):** Get a (learned) embedding for each token from the embedding table.
- **Word Positional Encodings (WPE):** Get a positional encoding for each position. Can be learned or fixed.
- Sum WPE and WTE for each token. This is the input to the transformer.
- **Apply transformer.**
- **Project onto vocabulary**
- **Apply softmax** to obtain N distribution(s) over next tokens, one for each time-step.
- **Train to minimize cross-entropy loss.**

GPT2-Style Decoder



```
N = len(idx)

# Sum the word embeddings and positional encodings.
wte = self.wte(idx)
wpe = self.wpe(jnp.arange(N))
x = wte + wpe

# Forward through the transformer.
x = self.transformer(x)

# Project onto the vocabulary.
x = jax.vmap(self.vproj)(x)
```